

Proximal Gradient Methods with Adaptive Subspace Sampling



Dmitry Grishchenko

MOTOR 2021
Online

5 July 2021

Collaborators



F. Iutzeler
LJK



J. Malick
CNRS, LJK



Dmitry Grishchenko, Franck Iutzeler, and Jérôme Malick. *Proximal gradient methods with adaptive subspace sampling*. Mathematics of Operations Research, 2021.

Outline



Motivation

Randomized Subspace Descent

Adaptive Randomized Subspace Descent

Outline



Motivation

Randomized Subspace Descent

Adaptive Randomized Subspace Descent



ML as an Optimization Problem

Empirical Risk Minimization

Loss function: represents the difference between two arguments.

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)}$$



ML as an Optimization Problem

Empirical Risk Minimization

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)}$$

Loss function: represents the difference between two arguments.

Learning is a compromise between accuracy and complexity

ML as an Optimization Problem

Structural Risk Minimization

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)} + r(x)$$

Loss function: represents the difference between two arguments.

Regularization penalty.



ML as an Optimization Problem

Structural Risk Minimization

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)} + r(x)$$

Smooth, convex.
Convex, non-smooth.

ML as an Optimization Problem

Structural Risk Minimization

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)} + r(x)$$

Smooth, convex.

Convex, non-smooth.

Why non smoothness?

ML as an Optimization Problem

Structural Risk Minimization

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x))}_{f(x)} + r(x)$$

Smooth, convex.
Convex, non-smooth.

Why non smoothness?

To enforce some structure of the optimal solution.

Sparse solution $r = \|\cdot\|_1$,

e.g. feature selection problems

Fixed variation $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

e.g. signal processing



Samuel Vaiter et al. *Model selection with low complexity priors*. Information and Inference: A Journal of the IMA 4.3 (2015): 230-287.



Proximal Gradient Descent

Let us consider a composite optimization problem

$$\min_{x \in \mathbb{R}^n} f(x) + r(x),$$

where f is L -smooth and convex, and r is convex, l.s.c.

Proximal Gradient Descent

Let us consider a composite optimization problem

$$\min_{x \in \mathbb{R}^n} f(x) + r(x),$$

where f is L -smooth and convex, and r is convex, l.s.c.

Proximal operator

$$\text{prox}_r(y) = \operatorname{argmin}_{x \in \mathbb{R}^n} \left\{ r(x) + \frac{1}{2} \|x - y\|_2^2 \right\}.$$

This operator is well defined for convex r and has a closed form solution for relatively simple r .

Proximal Gradient Descent

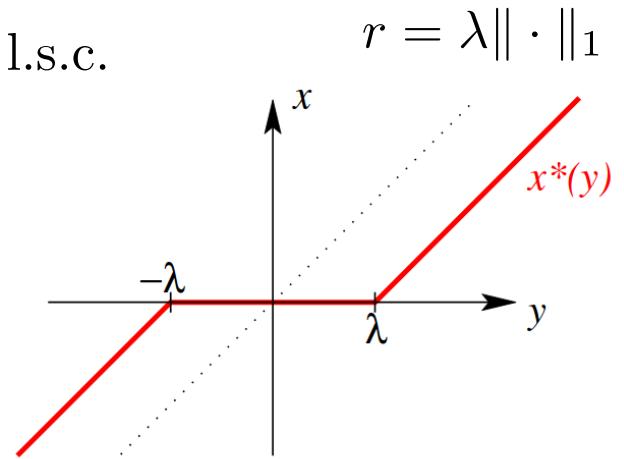
Let us consider a composite optimization problem

$$\min_{x \in \mathbb{R}^n} f(x) + r(x),$$

where f is L -smooth and convex, and r is convex, l.s.c.

Proximal operator

$$\text{prox}_r(y) = \operatorname{argmin}_{x \in \mathbb{R}^n} \left\{ r(x) + \frac{1}{2} \|x - y\|_2^2 \right\}.$$



This operator is well defined for convex r and has a closed form solution for relatively simple r .

Proximal Gradient Descent

Let us consider a composite optimization problem

$$\min_{x \in \mathbb{R}^n} f(x) + r(x),$$

where f is L -smooth and convex, and r is convex, l.s.c.

Proximal gradient descent

Step 1 $y^k = x^k - \gamma \nabla f(x)$ **forward (gradient) step.**

Step 2 $x^{k+1} = \text{prox}_{\gamma r}(y^k)$ **backward (proximal) step.**



R Tyrrell Rockafellar. *Monotone operators and the proximal point algorithm.*
SIAM journal on control and optimization, 14(5):877–898, 1976.

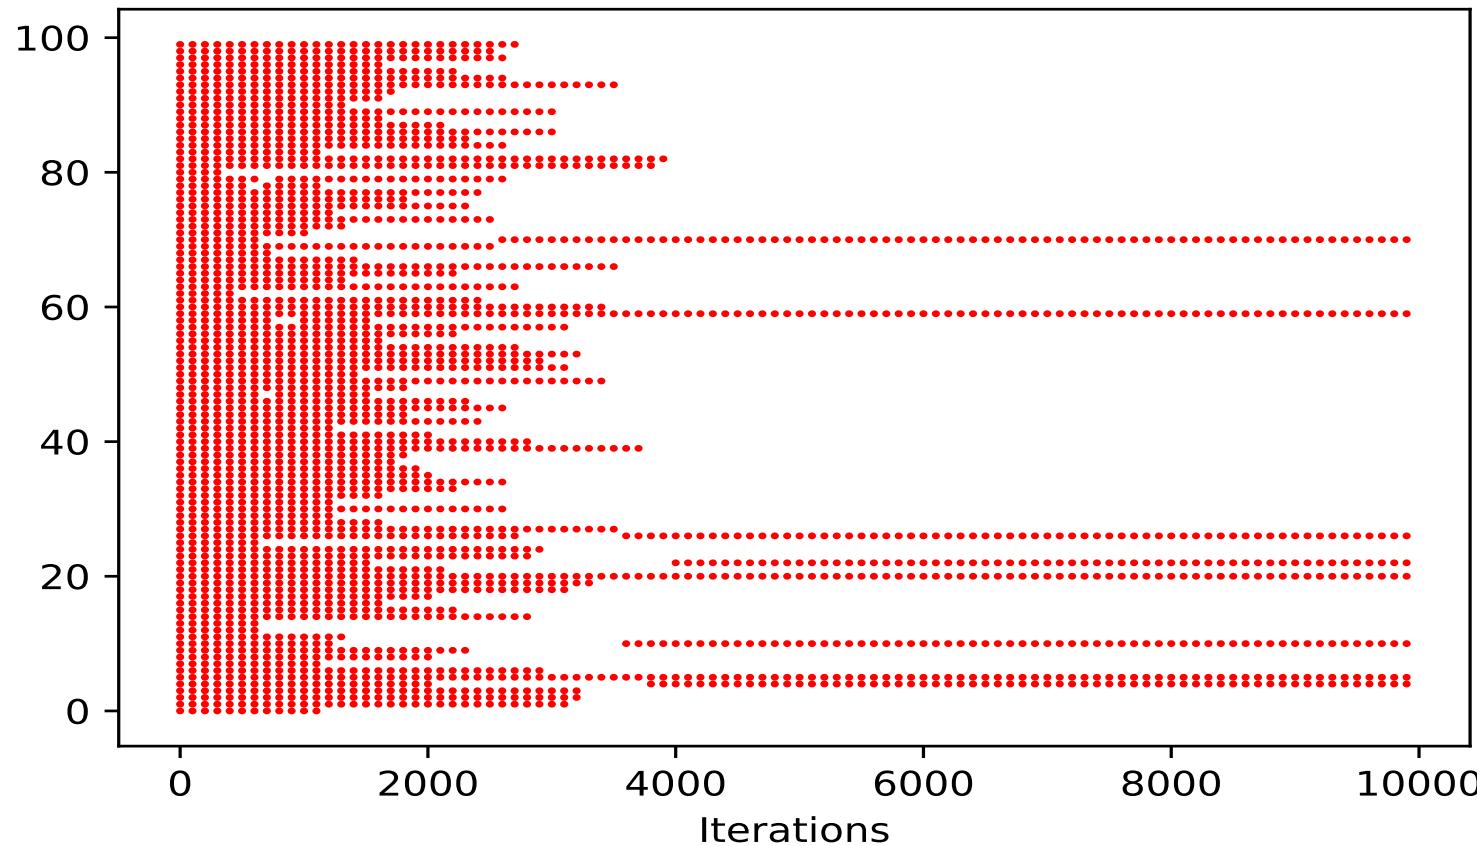
Identification

One nice thing

Proximal methods identify a near optimal subspace.



Coordinates in the support



Synthetic LASSO problem $\min \frac{1}{2} \|Ax - b\|_2^2 + \lambda_1 \|x\|_1$ for random generated matrix $A \in \mathbb{R}^{100 \times 100}$ and vector $b \in \mathbb{R}^{100}$ and hyperparameter λ_1 chosen to reach 8% of density (amount of non-zero coordinates) of the final solution.

Identification

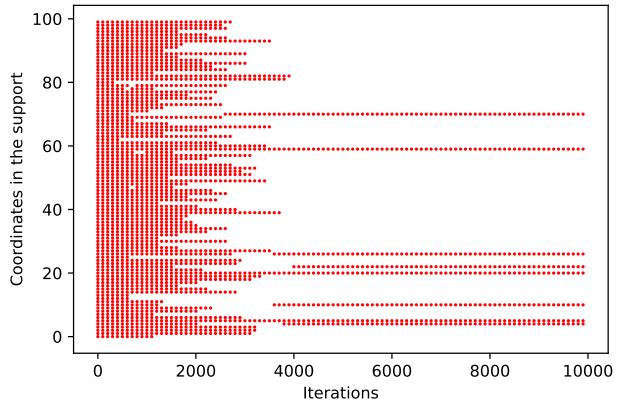
One nice thing

Proximal methods identify a near optimal subspace.

Sparsity vector

Let $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_m\}$ be a family of subspaces of \mathbb{R}^n with m elements. We define the sparsity vector on \mathcal{M} for point $x \in \mathbb{R}^n$ as the $\{0, 1\}$ -valued vector $S_{\mathcal{M}}(x) \in \{0, 1\}^m$ verifying

$$(S_{\mathcal{M}}(x))_{[i]} = 0 \quad \text{if } x \in \mathcal{M}_i \text{ and } 1 \text{ elsewhere.}$$



Identification

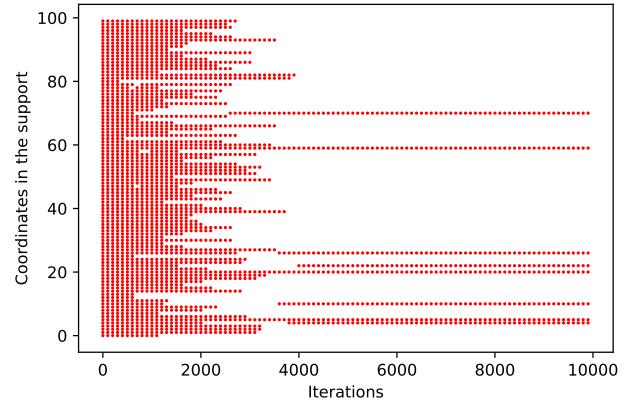
One nice thing

Proximal methods identify a near optimal subspace.

Sparsity vector

Let $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_m\}$ be a family of subspaces of \mathbb{R}^n with m elements. We define the sparsity vector on \mathcal{M} for point $x \in \mathbb{R}^n$ as the $\{0, 1\}$ -valued vector $S_{\mathcal{M}}(x) \in \{0, 1\}^m$ verifying

$$(S_{\mathcal{M}}(x))_{[i]} = 0 \quad \text{if } x \in \mathcal{M}_i \text{ and } 1 \text{ elsewhere.}$$



The collection $\mathcal{M} = \{\mathcal{M}_i\}_{1 \leq i \leq n}$ is the set of subspaces \mathcal{M}_i with $\text{supp}(x) = [n] \setminus \{i\}$ for all $x \in \mathcal{M}_i$.

Identification

$$x^* = \operatorname{argmin}_{x \in \mathbb{R}^n} f(x) + r(x)$$



One nice thing

Proximal methods identify a near optimal subspace.

Sparsity vector

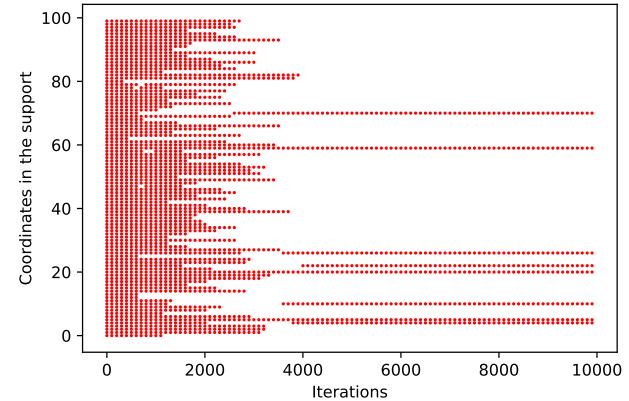
Let $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_m\}$ be a family of subspaces of \mathbb{R}^n with m elements. We define the sparsity vector on \mathcal{M} for point $x \in \mathbb{R}^n$ as the $\{0, 1\}$ -valued vector $S_{\mathcal{M}}(x) \in \{0, 1\}^m$ verifying

$$(S_{\mathcal{M}}(x))_{[i]} = 0 \quad \text{if } x \in \mathcal{M}_i \text{ and } 1 \text{ elsewhere.}$$

Theorem (Enlarged identification)

Let (u^k) be an \mathbb{R}^n -valued sequence converging almost surely to u^* and define sequence (x^k) as $x^k = \operatorname{prox}_{\gamma r}(u^k)$ and $x^* = \operatorname{prox}_{\gamma r}(u^*)$. Then (x^k) identifies some subspaces with probability one; more precisely for any $\varepsilon > 0$, with probability one, after some finite time,

$$S_{\mathcal{M}}(x^*) \leq S_{\mathcal{M}}(x^k) \leq \max \{S_{\mathcal{M}}(\operatorname{prox}_{\gamma r}(u)): u \in \mathcal{B}(u^*, \varepsilon)\}.$$



The collection $\mathcal{M} = \{\mathcal{M}_i\}_{1 \leq i \leq n}$ is the set of subspaces \mathcal{M}_i with $\operatorname{supp}(x) = [n] \setminus \{i\}$ for all $x \in \mathcal{M}_i$.



Franck Lutzeler and Jérôme Malick. *Nonsmoothness in Machine Learning: specific structure, proximal identification, and applications*. Set-Valued and Variational Analysis (2020): 1-18.

Identification

$$x^* = \operatorname{argmin}_{x \in \mathbb{R}^n} f(x) + r(x)$$



One nice thing

Proximal methods identify a near optimal subspace.

Sparsity vector

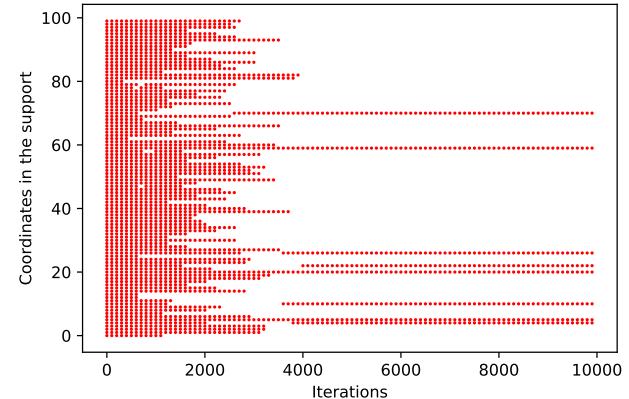
Let $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_m\}$ be a family of subspaces of \mathbb{R}^n with m elements. We define the sparsity vector on \mathcal{M} for point $x \in \mathbb{R}^n$ as the $\{0, 1\}$ -valued vector $S_{\mathcal{M}}(x) \in \{0, 1\}^m$ verifying

$$(S_{\mathcal{M}}(x))_{[i]} = 0 \quad \text{if } x \in \mathcal{M}_i \text{ and } 1 \text{ elsewhere.}$$

Theorem (Enlarged identification)

Let (u^k) be an \mathbb{R}^n -valued sequence converging almost surely to u^* and define sequence (x^k) as $x^k = \operatorname{prox}_{\gamma r}(u^k)$ and $x^* = \operatorname{prox}_{\gamma r}(u^*)$. Then (x^k) identifies some subspaces with probability one; more precisely for any $\varepsilon > 0$, with probability one, after some finite time,

$$S_{\mathcal{M}}(x^*) \leq S_{\mathcal{M}}(x^k) \leq \max \{S_{\mathcal{M}}(\operatorname{prox}_{\gamma r}(u)): u \in \mathcal{B}(u^*, \varepsilon)\}.$$



The collection $\mathcal{M} = \{\mathcal{M}_i\}_{1 \leq i \leq n}$ is the set of subspaces \mathcal{M}_i with $\operatorname{supp}(x) = [n] \setminus \{i\}$ for all $x \in \mathcal{M}_i$.

$$\begin{aligned} \operatorname{supp}(x^*) &\subseteq \operatorname{supp}(x^k) \\ &\subseteq \max_{u \in \mathcal{B}(u^*, \varepsilon)} \{\operatorname{supp}(\operatorname{prox}_{\gamma r}(u))\}. \end{aligned}$$



Franck Lutzeler and Jérôme Malick. *Nonsmoothness in Machine Learning: specific structure, proximal identification, and applications*. Set-Valued and Variational Analysis (2020): 1-18.

Identification

$$x^* = \operatorname{argmin}_{x \in \mathbb{R}^n} f(x) + r(x)$$



One nice thing

Proximal methods identify a near optimal subspace.

Sparsity vector

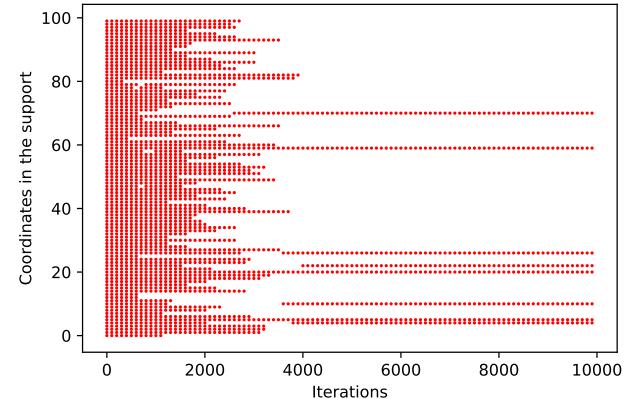
Let $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_m\}$ be a family of subspaces of \mathbb{R}^n with m elements. We define the sparsity vector on \mathcal{M} for point $x \in \mathbb{R}^n$ as the $\{0, 1\}$ -valued vector $S_{\mathcal{M}}(x) \in \{0, 1\}^m$ verifying

$$(S_{\mathcal{M}}(x))_{[i]} = 0 \quad \text{if } x \in \mathcal{M}_i \text{ and } 1 \text{ elsewhere.}$$

Theorem (Enlarged identification)

Let (u^k) be an \mathbb{R}^n -valued sequence converging almost surely to u^* and define sequence (x^k) as $x^k = \operatorname{prox}_{\gamma r}(u^k)$ and $x^* = \operatorname{prox}_{\gamma r}(u^*)$. Then (x^k) identifies some subspaces with probability one; more precisely for any $\varepsilon > 0$, with probability one, after some finite time,

$$S_{\mathcal{M}}(x^*) \leq S_{\mathcal{M}}(x^k) \leq \max \{S_{\mathcal{M}}(\operatorname{prox}_{\gamma r}(u)): u \in \mathcal{B}(u^*, \varepsilon)\}.$$



The collection $\mathcal{M} = \{\mathcal{M}_i\}_{1 \leq i \leq n}$ is the set of subspaces \mathcal{M}_i with $\operatorname{supp}(x) = [n] \setminus \{i\}$ for all $x \in \mathcal{M}_i$.

$$\begin{aligned} \operatorname{supp}(x^*) &\subseteq \operatorname{supp}(x^k) \\ &\subseteq \max_{u \in \mathcal{B}(u^*, \varepsilon)} \{\operatorname{supp}(\operatorname{prox}_{\gamma r}(u))\}. \end{aligned}$$

The same = verifies QC.



Franck Lutzeler and Jérôme Malick. Nonsmoothness in Machine Learning: specific structure, proximal identification, and applications. Set-Valued and Variational Analysis (2020): 1-18.

Outline



Motivation

Randomized Subspace Descent

Adaptive Randomized Subspace Descent

Randomized Coordinate Descent



Full gradient computation is expensive.

Randomized Coordinate Descent



Full gradient computation is expensive.

Coordinate descent methods is a class of iterative methods in which only one coordinate (block) is updated on every iteration.



Randomized Coordinate Descent

Full gradient computation is expensive.

Coordinate descent methods is a class of iterative methods in which only one coordinate (block) is updated on every iteration.

Example 1 (smooth).

$$x^{k+1} = x^k - \gamma \nabla f(x)_{[i^k]}$$



Randomized Coordinate Descent

Full gradient computation is expensive.

Coordinate descent methods is a class of iterative methods in which only one coordinate (block) is updated on every iteration.

Example 1 (smooth).

$$x^{k+1} = x^k - \gamma \nabla f(x)_{[i^k]}$$

Example 2 (separable regularizer).

$$r(x) = \sum_{i=1}^n r_i(x_{[i]}) \Rightarrow \text{prox}_{\gamma r}(x)_{[i]} = \text{prox}_{\gamma r_i}(x_{[i]}).$$
$$x_{[i^k]}^{k+1} \leftarrow \text{prox}_{\gamma r_{i^k}} \left(x_{[i^k]}^k - \gamma \nabla_{[i^k]} f(x^k) \right)$$



Randomized Coordinate Descent

Full gradient computation is expensive.

Coordinate descent methods is a class of iterative methods in which only one coordinate (block) is updated on every iteration.

Example 1 (smooth).

$$x^{k+1} = x^k - \gamma \nabla f(x)_{[i^k]}$$

Example 2 (separable regularizer).

$$r(x) = \sum_{i=1}^n r_i(x_{[i]}) \Rightarrow \text{prox}_{\gamma r}(x)_{[i]} = \text{prox}_{\gamma r_i}(x_{[i]}).$$
$$x_{[i^k]}^{k+1} \leftarrow \text{prox}_{\gamma r_{i^k}} \left(x_{[i^k]}^k - \gamma \nabla_{[i^k]} f(x^k) \right)$$

Drawback: explicit use of the separability of the regularizer.



Peter Richtárik and Martin Takáč. *Iteration complexity of randomized block-coordinate descent methods for minimizing a composite function.* Mathematical Programming 144.1-2 (2014): 1-38.



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.



Olivier Fercoq and Pascal Bianchi. *A coordinate-descent primal-dual algorithm with large step size and possibly nonseparable functions*. SIAM Journal on Optimization 29.1 (2019): 100-134.



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x_{[i^k]}^{k+1} \leftarrow \text{prox}_{\gamma r_{i^k}} \left(x_{[i^k]}^k - \gamma \nabla_{[i^k]} f(x^k) \right)$$



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x_{[i^k]}^{k+1} \leftarrow \text{prox}_{\gamma r_{i^k}} \left(x_{[i^k]}^k - \gamma \nabla_{[i^k]} f(x^k) \right)$$



Randomized Subspace Descent

What if the regularizer is not separable? e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r_{i^k}} \left(x_{[i^k]}^k - \gamma \nabla_{[i^k]} f(x^k) \right) + [x^k]_{\bar{i}^k}$$



Randomized Subspace Descent

What if the regularizer is not separable? e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left([y^k]_{i^k} + [y^{k-1}]_{\bar{i}^k} \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left([y^k]_{i^k} + [y^{k-1}]_{\bar{i}^k} \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.

In this reformulation the separability is not required!



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left([y^k]_{i^k} + [y^{k-1}]_{\bar{i}^k} \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.

Two orthogonal projections
onto orthogonal spaces!

In this reformulation the separability is not required!



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left(P(y^k) + (I - P)(y^{k-1}) \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.

In this reformulation the separability is not required!



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left(P(y^k) + (I - P)(y^{k-1}) \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.

In this reformulation the separability is not required!

General orthogonal projections are used!



Randomized Subspace Descent

What if the regularizer is not separable?

e.g. $r = \sum_{i=1}^{n-1} |x_{i+1} - x_i|$.

$$x^{k+1} = \text{prox}_{\gamma r} \left(P(y^k) + (I - P)(y^{k-1}) \right),$$

where $y^k = x^k - \gamma \nabla f(x^k)$.

In this reformulation the separability is not required!

General orthogonal projections are used!

Does it work like this?

Examples: Subspaces





Examples: Subspaces

Example 3.

Let us consider the set of subspaces \mathcal{C}_i such that \mathcal{C}_i is i -th coordinate line. Select an orthogonal projection onto the \mathcal{C}_i with probability $\frac{1}{n-1}$ $\forall i \in [2, n]$ and 0 for the 1-st.



Examples: Subspaces

Example 3.

Let us consider the set of subspaces \mathcal{C}_i such that \mathcal{C}_i is i -th coordinate line. Select an orthogonal projection onto the \mathcal{C}_i with probability $\frac{1}{n-1} \forall i \in [2, n]$ and 0 for the 1-st.

Does not work if the first coordinates of the starting and the optimal point are different.

Examples: Subspaces

Example 3.

Let us consider the set of subspaces \mathcal{C}_i such that \mathcal{C}_i is i -th coordinate line. Select an orthogonal projection onto the \mathcal{C}_i with probability $\frac{1}{n-1} \forall i \in [2, n]$ and 0 for the 1-st.

Does not work if the first coordinates of the starting and the optimal point are different.

Covering family of subspaces

Let $\mathcal{C} = \{\mathcal{C}_i\}_i$ be a family of subspaces of \mathbb{R}^n . We say that \mathcal{C} is *covering* if it spans the whole space, i.e. if $\sum_i \mathcal{C}_i = \mathbb{R}^n$.

Admissible Selection





Admissible Selection

Let \mathcal{C} be a covering family of subspaces of \mathbb{R}^n . A selection \mathfrak{S} is defined from the set of all subsets of \mathcal{C} to the set of the subspaces of \mathbb{R}^n as

$$\mathfrak{S}(\omega) = \sum_{j=1}^s \mathcal{C}_{i_j} \quad \text{for } \omega = \{\mathcal{C}_{i_1}, \dots, \mathcal{C}_{i_s}\}.$$

The selection \mathfrak{S} is *admissible* if $\mathbb{P}[x \in \mathfrak{S}^\perp] < 1$ for all $x \in \mathbb{R}^n \setminus \{0\}$.



Admissible Selection

Let \mathcal{C} be a covering family of subspaces of \mathbb{R}^n . A selection \mathfrak{S} is defined from the set of all subsets of \mathcal{C} to the set of the subspaces of \mathbb{R}^n as

$$\mathfrak{S}(\omega) = \sum_{j=1}^s \mathcal{C}_{i_j} \quad \text{for } \omega = \{\mathcal{C}_{i_1}, \dots, \mathcal{C}_{i_s}\}.$$

The selection \mathfrak{S} is *admissible* if $\mathbb{P}[x \in \mathfrak{S}^\perp] < 1$ for all $x \in \mathbb{R}^n \setminus \{0\}$.

If a selection \mathfrak{S} is admissible then $\mathsf{P} := \mathbb{E}[P_{\mathfrak{S}}]$ is a positive definite matrix.

In this case, we denote by $\lambda_{\min}(\mathsf{P}) > 0$ and $\lambda_{\max}(\mathsf{P}) \leq 1$ its minimal and maximal eigenvalues.

Algorithm 1: RPSD

Algorithm 1 Randomized Proximal Subspace Descent - RPSD

- 1: Input: $Q = P^{-\frac{1}{2}}$
 - 2: Initialize z^0 , $x^1 = \text{prox}_{\gamma r}(Q^{-1}(z^0))$
 - 3: **for** $k = 1, \dots$ **do**
 - 4: $y^k = Q(x^k - \gamma \nabla f(x^k))$
 - 5: $z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1})$
 - 6: $x^{k+1} = \text{prox}_{\gamma r}(Q^{-1}(z^k))$
 - 7: **end for**
-

Algorithm 1: RPSD

Algorithm 1 Randomized Proximal Subspace Descent - RPSD

```

1: Input:  $Q = P^{-\frac{1}{2}}$ 
2: Initialize  $z^0$ ,  $x^1 = \text{prox}_{\gamma r}(Q^{-1}(z^0))$ 
3: for  $k = 1, \dots$  do
4:    $y^k = Q(x^k - \gamma \nabla f(x^k))$  ← “Sketch”
5:    $z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1})$  ← “Project”
6:    $x^{k+1} = \text{prox}_{\gamma r}(Q^{-1}(z^k))$ 
7: end for

```

RPSD: Convergence Result



Assumption (on randomness)

Given a covering family $\mathcal{C} = \{\mathcal{C}_i\}$ of subspaces, we consider a sequence $\mathfrak{S}^1, \mathfrak{S}^2, \dots, \mathfrak{S}^k$ of admissible selections, which is i.i.d.

RPSD: Convergence Result

Assumption (on randomness)

Given a covering family $\mathcal{C} = \{\mathcal{C}_i\}$ of subspaces, we consider a sequence $\mathfrak{S}^1, \mathfrak{S}^2, \dots, \mathfrak{S}^k$ of admissible selections, which is i.i.d.

Theorem (Convergence of RPSD)

For any $\gamma \in (0, 2/(\mu+L)]$, the sequence (x^k) of the iterates of RPSD converges almost surely to the minimizer x^* with rate

$$\mathbb{E} [\|x^{k+1} - x^*\|_2^2] \leq \left(1 - \lambda_{\min}(\mathsf{P}) \frac{2\gamma\mu L}{\mu + L}\right)^k C,$$

where $C = \lambda_{\max}(\mathsf{P}) \|z^0 - \mathsf{Q}(x^* - \gamma \nabla f(x^*))\|_2^2$.

RPSD: Convergence Result

Consider the set of subspaces \mathcal{C}_i such that \mathcal{C}_i is i -th coordinate line. Consider the selection \mathfrak{S} such that $\mathbb{P}[\mathcal{C}_i \in \mathfrak{S}] = p_i > 0$, then $\lambda_{\min}(\mathsf{P}) = \min_i p_i > 0$.

Theorem (Convergence of RPSD)

For any $\gamma \in (0, 2/(\mu+L)]$, the sequence (x^k) of the iterates of RPSD converges almost surely to the minimizer x^* with rate

$$\mathbb{E} [\|x^{k+1} - x^*\|_2^2] \leq \left(1 - \lambda_{\min}(\mathsf{P}) \frac{2\gamma\mu L}{\mu + L}\right)^k C,$$

where $C = \lambda_{\max}(\mathsf{P}) \|z^0 - \mathsf{Q}(x^* - \gamma \nabla f(x^*))\|_2^2$.



RPSD: Proof Sketch

Lemma 1

From the minimizer x^* , define the fixed points $z^* = y^* = \mathbf{Q}(x^* - \gamma \nabla f(x^*))$ of the sequences (y^k) and (z^k) . Then

$$\mathbb{E} [\|z^k - z^*\|_2^2 | \mathcal{F}^{k-1}] = \|z^{k-1} - z^*\|_2^2 + \|y^k - y^*\|_{\mathbf{P}}^2 - \|z^{k-1} - z^*\|_{\mathbf{P}}^2,$$

where $\mathcal{F}^k = \sigma(\{\mathfrak{S}_\ell\}_{\ell \leq k})$ is the filtration of the past random subspaces.

$$z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1})$$

RPSD: Proof Sketch

Lemma 1

From the minimizer x^* , define the fixed points $z^* = y^* = \mathbf{Q}(x^* - \gamma \nabla f(x^*))$ of the sequences (y^k) and (z^k) . Then

$$\mathbb{E} [\|z^k - z^*\|_2^2 | \mathcal{F}^{k-1}] = \|z^{k-1} - z^*\|_2^2 + \|y^k - y^*\|_{\mathsf{P}}^2 - \|z^{k-1} - z^*\|_{\mathsf{P}}^2,$$

where $\mathcal{F}^k = \sigma(\{\mathfrak{S}_\ell\}_{\ell \leq k})$ is the filtration of the past random subspaces.

Lemma 2

Using the same notations as in Lemma 1

$$\|y^k - y^*\|_{\mathsf{P}}^2 - \|z^{k-1} - z^*\|_{\mathsf{P}}^2 \leq -\lambda_{\min}(\mathsf{P}) \frac{2\gamma\mu L}{\mu + L} \|z^{k-1} - z^*\|_2^2.$$

RPSD: Proof Sketch

Lemma 1

From the minimizer x^* , define the fixed points $z^* = y^* = \mathbf{Q}(x^* - \gamma \nabla f(x^*))$ of the sequences (y^k) and (z^k) . Then

$$\mathbb{E} [\|z^k - z^*\|_2^2 | \mathcal{F}^{k-1}] = \|z^{k-1} - z^*\|_2^2 + \|y^k - y^*\|_{\mathbf{P}}^2 - \|z^{k-1} - z^*\|_{\mathbf{P}}^2,$$

where $\mathcal{F}^k = \sigma(\{\mathfrak{S}_\ell\}_{\ell \leq k})$ is the filtration of the past random subspaces.

Lemma 2

Using the same notations as in Lemma 1



$$\|y^k - y^*\|_{\mathbf{P}}^2 - \|z^{k-1} - z^*\|_{\mathbf{P}}^2 \leq -\lambda_{\min}(\mathbf{P}) \frac{2\gamma\mu L}{\mu + L} \|z^{k-1} - z^*\|_2^2.$$



Examples: TV Projections

$$r = \lambda \sum_{i=1}^{n-1} |x_{[i]} - x_{[i+1]}|$$

Fixed variation sparsity = small amount of blocks of equal coordinates.



Examples: TV Projections

$$r = \lambda \sum_{i=1}^{n-1} |x_{[i]} - x_{[i+1]}|$$

Fixed variation sparsity = small amount of blocks of equal coordinates.

Projection on such set

$$P_{\mathfrak{S}} = \left(\begin{array}{ccccccccc} \overbrace{\frac{1}{n_1} & \cdots & \frac{1}{n_1}}^{n_1} & 0 & \cdots & \overbrace{\cdots & \cdots}^{n-n_s} & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots & \vdots \\ \frac{1}{n_1} & \cdots & \frac{1}{n_1} & 0 & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & \cdots & 0 & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & 0 & \cdots & 0 & \\ \vdots & \ddots & \ddots & \ddots & 0 & \frac{1}{n-n_s} & \cdots & \frac{1}{n-n_s} & \\ \vdots & \ddots & \ddots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & \cdots & \cdots & \cdots & 0 & \frac{1}{n-n_s} & \cdots & \frac{1}{n-n_s} & \end{array} \right)_{n-n_s}$$

Outline



Motivation

Randomized Subspace Descent

Adaptive Randomized Subspace Descent

Algorithm 2: ARPSD

Algorithm 2 Adaptive Randomized Proximal Subspace Descent - ARPSD

Initialize z^0 , $x^1 = \text{prox}_{\gamma g}(Q_0^{-1}(z^0))$, $\ell = 0$, $L = \{0\}$.

for $k = 1, \dots$ **do**

$$y^k = Q_\ell(x^k - \gamma \nabla f(x^k))$$

$$z^k = P_{S^k}(y^k) + (I - P_{S^k})(z^{k-1})$$

$$x^{k+1} = \text{prox}_{\gamma g}(Q_\ell^{-1}(z^k))$$

if an adaptation is decided **then**

$$L \leftarrow L \cup \{k + 1\}, \ell \leftarrow \ell + 1$$

Generate a new admissible selection

Compute $Q_\ell = P_\ell^{-\frac{1}{2}}$ and Q_ℓ^{-1}

Rescale $z^k \leftarrow Q_\ell Q_{\ell-1}^{-1} z^k$

end if

end for

Algorithm 2: ARPSD

Algorithm 2 Adaptive Randomized Proximal Subspace Descent - ARPSD

Initialize z^0 , $x^1 = \text{prox}_{\gamma g}(Q_0^{-1}(z^0))$, $\ell = 0$, $L = \{0\}$.

for $k = 1, \dots$ **do**

$$y^k = Q_\ell(x^k - \gamma \nabla f(x^k))$$

$$z^k = P_{S^k}(y^k) + (I - P_{S^k})(z^{k-1})$$

$$x^{k+1} = \text{prox}_{\gamma g}(Q_\ell^{-1}(z^k))$$

if an adaptation is decided **then**

$$L \leftarrow L \cup \{k + 1\}, \ell \leftarrow \ell + 1$$

Generate a new admissible selection

Compute $Q_\ell = P_\ell^{-\frac{1}{2}}$ and Q_ℓ^{-1}

Rescale $z^k \leftarrow Q_\ell Q_{\ell-1}^{-1} z^k$

end if

end for

Algorithm 2: ARPSD

Algorithm 2 Adaptive Randomized Proximal Subspace Descent - ARPSD

Initialize z^0 , $x^1 = \text{prox}_{\gamma g}(\mathbf{Q}_0^{-1}(z^0))$, $\ell = 0$, $\mathcal{L} = \{0\}$.

for $k = 1, \dots$ **do**

$$y^k = \mathbf{Q}_\ell(x^k - \gamma \nabla f(x^k))$$

$$z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1})$$

$$x^{k+1} = \text{prox}_{\gamma g}(\mathbf{Q}_\ell^{-1}(z^k))$$

if an adaptation is decided **then**

$$\mathcal{L} \leftarrow \mathcal{L} \cup \{k+1\}, \ell \leftarrow \ell + 1$$

Generate a new admissible selection

Compute $\mathbf{Q}_\ell = \mathbf{P}_\ell^{-\frac{1}{2}}$ and \mathbf{Q}_ℓ^{-1}

Rescale $z^k \leftarrow \mathbf{Q}_\ell \mathbf{Q}_{\ell-1}^{-1} z^k$

end if

end for

$$\mathbb{P}[\mathcal{C}_i \in \mathfrak{S}^{k+1}] = \begin{cases} p & \text{if } x^{k+1} \in \mathcal{M}_i \Leftrightarrow [\mathbf{S}_{\mathcal{M}}(x^{k+1})]_i = 0 \\ 1 & \text{elsewhere} \end{cases}$$

Algorithm 2: ARPSD

Algorithm 2 Adaptive Randomized Proximal Subspace Descent - ARPSD

Initialize z^0 , $x^1 = \text{prox}_{\gamma g}(\mathbf{Q}_0^{-1}(z^0))$, $\ell = 0$, $\mathcal{L} = \{0\}$.

for $k = 1, \dots$ **do**

$$y^k = \mathbf{Q}_\ell(x^k - \gamma \nabla f(x^k))$$

$$z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1})$$

$$x^{k+1} = \text{prox}_{\gamma g}(\mathbf{Q}_\ell^{-1}(z^k))$$

if an adaptation is decided **then**

$$\mathcal{L} \leftarrow \mathcal{L} \cup \{k + 1\}, \ell \leftarrow \ell + 1$$

Generate a new admissible selection

Compute $\mathbf{Q}_\ell = \mathbf{P}_\ell^{-\frac{1}{2}}$ and \mathbf{Q}_ℓ^{-1}

$$\text{Rescale } z^k \leftarrow \mathbf{Q}_\ell \mathbf{Q}_{\ell-1}^{-1} z^k$$

end if

end for



Adaptation Process

Let us specify ARPSD with the following simple adaptation strategy. We take a fixed upper bound on the adaptation cost and a fixed lower bound on uniformity:

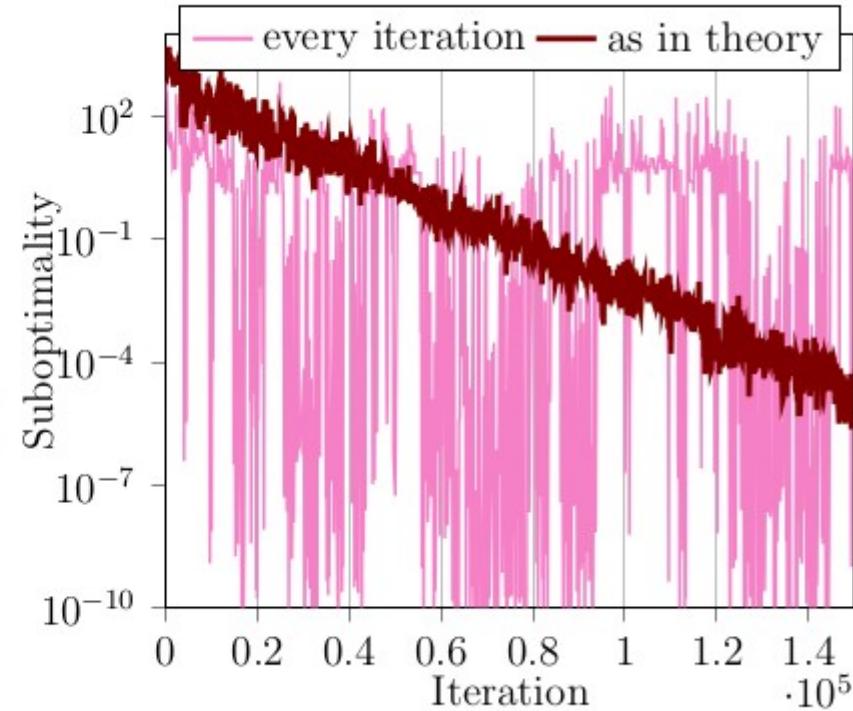
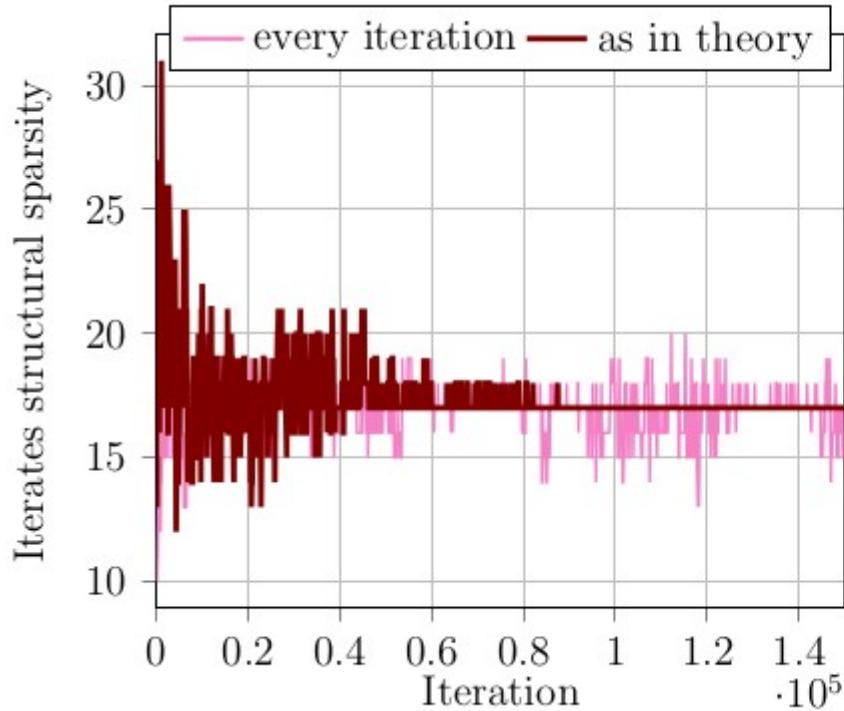
$$\|Q_\ell Q_{\ell-1}^{-1}\|_2^2 \leq a \quad \lambda_{\min}(P_\ell) \geq \lambda.$$

Then from the rate $1 - \alpha = 1 - 2\gamma\mu L\lambda/(\mu + L)$, we can perform an adaptation every

$$c = \lceil \log(a) / \log((2 - \alpha)/(2 - 2\alpha)) \rceil$$

iterations, so that $a(1 - \alpha)^c = (1 - \alpha/2)^c$ and $k_\ell = \ell c$.

Adaptation Process



ARPSD: Convergence Result



ARPSD: Convergence Result

Assumption (on randomness)

For all $k > 0$, \mathfrak{S}^k is \mathcal{F}^k -measurable and admissible. Furthermore, if $k \notin L$, (\mathfrak{S}^k) is independent and identically distributed on $[k_\ell, k]$. The decision to adapt or not at time k is \mathcal{F}^k -measurable, i.e. $(k_\ell)_\ell$ is a sequence of \mathcal{F}^k -stopping times.

ARPSD: Convergence Result

Assumption (on randomness)

For all $k > 0$, \mathfrak{S}^k is \mathcal{F}^k -measurable and admissible. Furthermore, if $k \notin \mathbb{L}$, (\mathfrak{S}^k) is independent and identically distributed on $[k_\ell, k]$. The decision to adapt or not at time k is \mathcal{F}^k -measurable, i.e. $(k_\ell)_\ell$ is a sequence of \mathcal{F}^k -stopping times.

Theorem (Convergence of ARPSD)

For any $\gamma \in (0, 2/(\mu + L)]$, the sequence (x^k) of the iterates of ARPSD converges almost surely to the minimizer x^* with rate

$$\mathbb{E} [\|x^{k+1} - x_\ell^*\|_2^2] \leq \left(1 - \frac{\lambda}{2} \frac{2\gamma\mu L}{\mu + L}\right)^k C.$$

where $C = \lambda_{\max}(\mathsf{P}) \|z^0 - \mathsf{Q}(x^* - \gamma \nabla f(x^*))\|_2^2$.

ARPSD: Convergence Result

Assumption (on randomness)

For all $k > 0$, \mathfrak{S}^k is \mathcal{F}^k -measurable and admissible. Furthermore, if $k \notin \mathbb{L}$, (\mathfrak{S}^k) is independent and identically distributed on $[k_\ell, k]$. The decision to adapt or not at time k is \mathcal{F}^k -measurable, i.e. $(k_\ell)_\ell$ is a sequence of \mathcal{F}^k -stopping times.

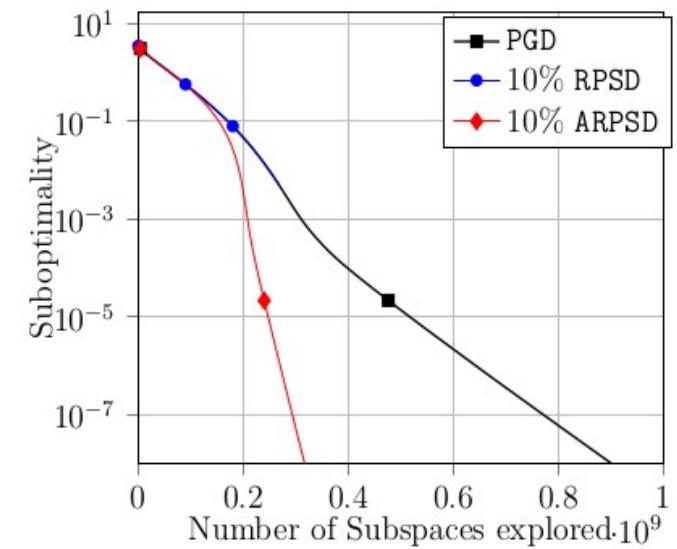
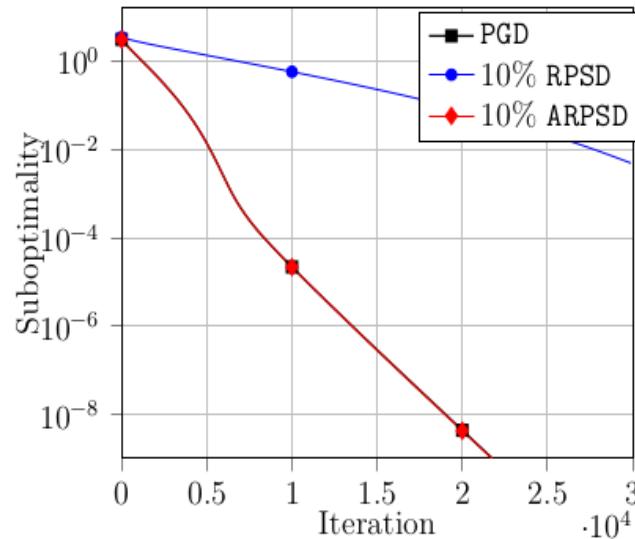
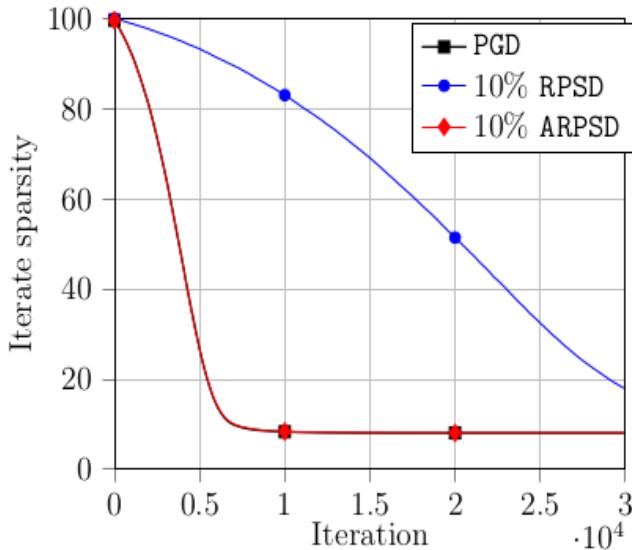
Theorem (Convergence of ARPSD)

For any $\gamma \in (0, 2/(\mu + L)]$, the sequence (x^k) of the iterates of ARPSD converges almost surely to the minimizer x^* with rate

$$\mathbb{E} [\|x^{k+1} - x_\ell^*\|_2^2] \leq \left(1 - \frac{\lambda}{2} \frac{2\gamma\mu L}{\mu + L}\right)^k C.$$

where $C = \lambda_{\max}(\mathsf{P}) \|z^0 - \mathsf{Q}(x^* - \gamma \nabla f(x^*))\|_2^2$.

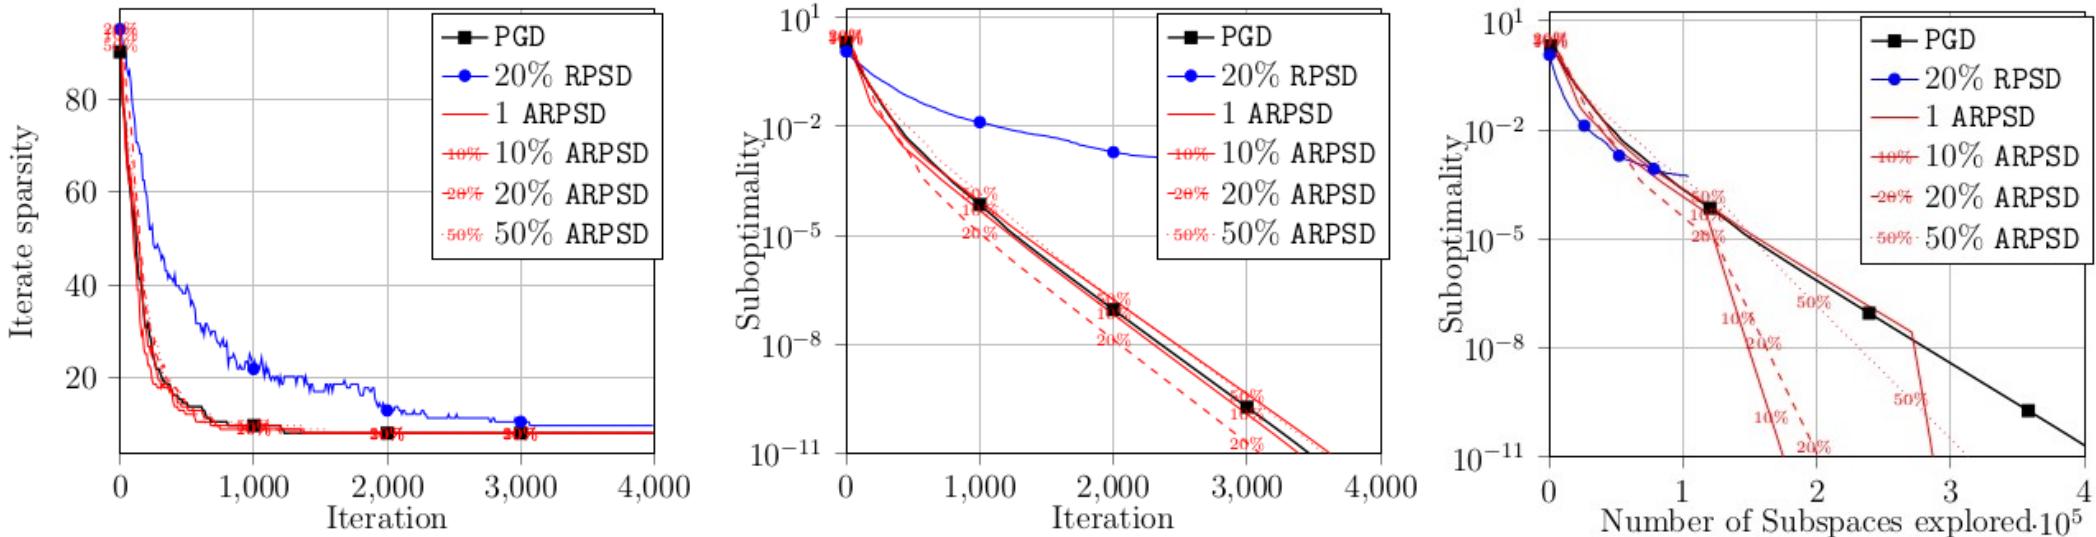
Experiments: Inefficiency of RPSD



Logistic regression with elastic net regularizer on rcv1-train dataset ($n = 47236$ $m = 20242$).

$$\min_{x \in \mathbb{R}^n} \frac{1}{m} \sum_{j=1}^m \log(1 + \exp(-y_j z_j^\top x)) + \lambda_1 \|x\|_1 + \frac{\lambda_2}{2} \|x\|_2^2$$

Experiments: ARPSD with TV



1D-TV-regularized logistic regression on a1a dataset ($n = 123$ $m = 1605$).

$$\min_{x \in \mathbb{R}^n} \frac{1}{m} \sum_{j=1}^m \log(1 + \exp(-y_j z_j^\top x)) + \lambda_1 \sum_{i=1}^{n-1} |x_{[i]} - x_{[i+1]}| + \frac{\lambda_2}{2} \|x\|_2^2$$



Strange Metric?

$$\min_{x \in \mathbb{R}^n} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(b_i, h(a_i, x)) + r(x)}_{f(x)}$$



Strange Metric?

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^M \alpha_i \underbrace{\left[\frac{1}{|\mathcal{D}_i|} \sum_{j \in \mathcal{D}_i} \ell(b_j, h(a_j, x)) \right]}_{f_i} + r(x),$$

where the full dataset \mathcal{D} is split onto M nonintersecting subsets \mathcal{D}_i and α_i is the proportion of examples $\frac{|\mathcal{D}_i|}{m}$.



Strange Metric?

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^M \alpha_i \underbrace{\left[\frac{1}{|\mathcal{D}_i|} \sum_{j \in \mathcal{D}_i} \ell(b_j, h(a_j, x)) \right]}_{f_i} + r(x),$$

where the full dataset \mathcal{D} is split onto M nonintersecting subsets \mathcal{D}_i and α_i is the proportion of examples $\frac{|\mathcal{D}_i|}{m}$.

These subsets \mathcal{D}_i can be split over machines.

Strange Metric?

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^M \alpha_i \underbrace{\left[\frac{1}{|\mathcal{D}_i|} \sum_{j \in \mathcal{D}_i} \ell(b_j, h(a_j, x)) \right]}_{f_i} + r(x),$$

where the full dataset \mathcal{D} is split onto M nonintersecting subsets \mathcal{D}_i and α_i is the proportion of examples $\frac{|\mathcal{D}_i|}{m}$.

These subsets \mathcal{D}_i can be split over machines.

$$z^k = \sum_i \alpha_i z_i^k$$



Master

$$z_i^k = P_{\mathfrak{S}^k} (y_i^k) + (I - P_{\mathfrak{S}^k}) (z_i^{k-1})$$

Strange Metric?

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^M \alpha_i \underbrace{\left[\frac{1}{|\mathcal{D}_i|} \sum_{j \in \mathcal{D}_i} \ell(b_j, h(a_j, x)) \right]}_{f_i} + r(x),$$

where the full dataset \mathcal{D} is split onto M nonintersecting subsets \mathcal{D}_i and α_i is the proportion of examples $\frac{|\mathcal{D}_i|}{m}$.

These subsets \mathcal{D}_i can be split over machines.

$$z^k = \sum_i \alpha_i z_i^k$$

Bottleneck





Master

$$z_i^k = P_{\mathfrak{S}^k} (y_i^k) + (I - P_{\mathfrak{S}^k}) (z_i^{k-1})$$

Strange Metric?

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^M \alpha_i \underbrace{\left[\frac{1}{|\mathcal{D}_i|} \sum_{j \in \mathcal{D}_i} \ell(b_j, h(a_j, x)) \right]}_{f_i} + r(x),$$

where the full dataset \mathcal{D} is split onto M nonintersecting subsets \mathcal{D}_i and α_i is the proportion of examples $\frac{|\mathcal{D}_i|}{m}$.

These subsets \mathcal{D}_i can be split over machines.

$$z^k = \sum_i \alpha_i z_i^k$$

Bottleneck



Amount of subspaces explored



Master

$$z_i^k = P_{\mathfrak{S}^k} (y_i^k) + (I - P_{\mathfrak{S}^k}) (z_i^{k-1})$$



Thank You For

Your Attention!



Practical for TV regularizer

Consider the set of artificial jumps $\mathcal{S} = \{n_1, n_2, \dots, n_{l-1}\}$ and denote by $\mathcal{R} = \{i \notin \mathcal{S} : [\mathbf{S}_{\mathcal{M}}(x^k)]_i = 0\}$ the set of possible random entries. Fix the amount of sampled elements s and sample “first” element \mathcal{R}_0 uniformly in $\mathcal{R} = \{\mathcal{R}_i\}_{1 \leq i \leq r}$. Select “first s ” elements starting from \mathcal{R}_f considering the cyclic structure of the list of elements ($\mathcal{R}_{r+1} = \mathcal{R}_1$).

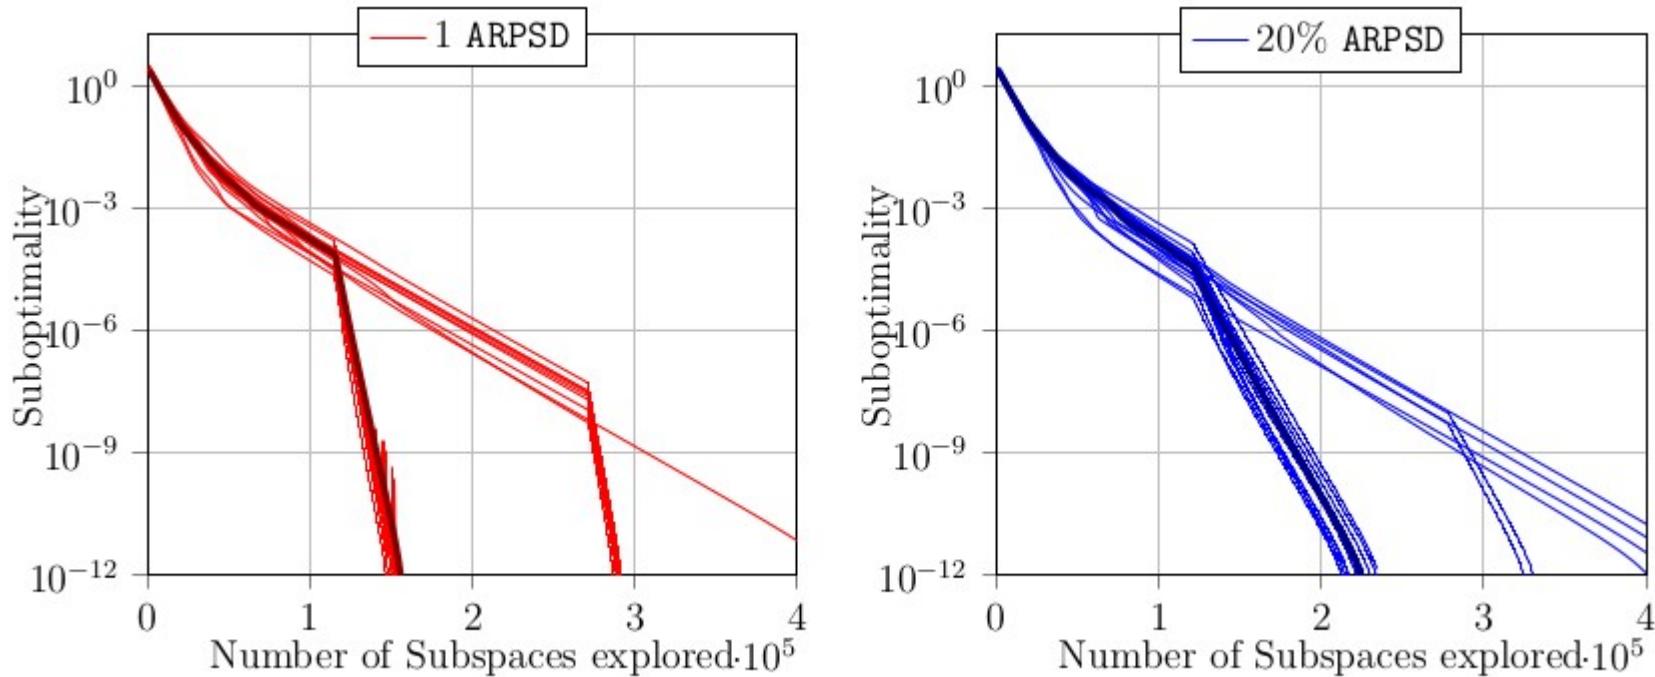
If l is small enough, it will not change the sparsity property of the random projection $P_{\mathcal{G}^k}$; however, this modification will force all the projections to be block-diagonal with blocks’ ends on positions n_1, \dots, n_{l-1} . In contrast with $\text{jumps}(x^k)$ that we could not control, by adding l artificial jumps, we could guarantee that each block of the $P_{\mathcal{G}^k}$ has at most $\lceil n/l \rceil$ rows. Since every random projection has end of the block on positions $\{n_i\}_{1 \leq i \leq l-1}$, \mathbf{P}_ℓ also has such block structure and we could split the computation of \mathbf{Q}_ℓ^{-1} and \mathbf{Q}_ℓ into l independent parts and could be done in parallel.

Strategies for (A)R PSD



	(non-adaptive) subspace descent RPSD	adaptive subspace descent ARPSD
Subspace family	$\mathcal{C} = \{\mathcal{C}_1, \dots, \mathcal{C}_c\}$	
Algorithm	$\begin{cases} y^k = Q(x^k - \gamma \nabla f(x^k)) \\ z^k = P_{\mathfrak{S}^k}(y^k) + (I - P_{\mathfrak{S}^k})(z^{k-1}) \\ x^{k+1} = \text{prox}_{\gamma g}(Q^{-1}(z^k)) \end{cases}$	
Selection	Option 1 $\mathcal{C}_i \in \mathfrak{S}^k$ with probability p	$\mathcal{C}_i \in \mathfrak{S}^k$ with probability $\begin{cases} p & \text{if } x^k \in \mathcal{M}_i \Leftrightarrow [\mathbf{S}_{\mathcal{M}}(x^k)]_i = 0 \\ 1 & \text{elsewhere} \end{cases}$
	Option 2 Sample s elements uniformly in uniformly in \mathcal{C}	Sample s elements uniformly in $\{\mathcal{C}_i : x^k \in \mathcal{M}_i \text{ i.e. } [\mathbf{S}_{\mathcal{M}}(x^k)]_i = 0\}$ and add <i>all</i> elements in $\{\mathcal{C}_j : x^k \notin \mathcal{M}_j \text{ i.e. } [\mathbf{S}_{\mathcal{M}}(x^k)]_j = 1\}$

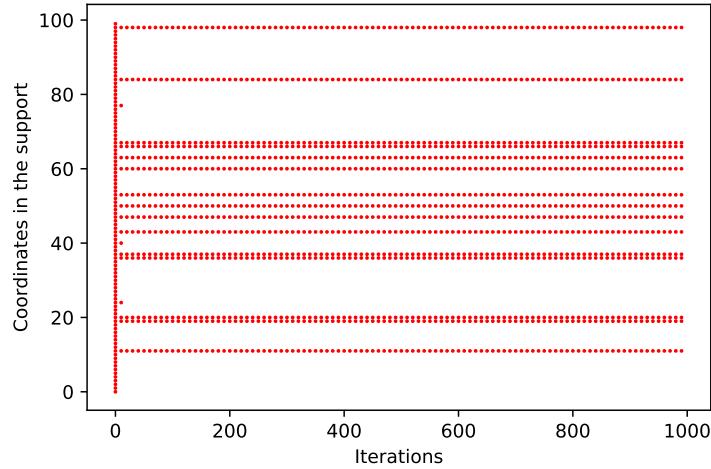
Practical robustness



Logistic regression with elastic net regularizer on rcv1_train dataset ($n = 47236$ $m = 20242$).

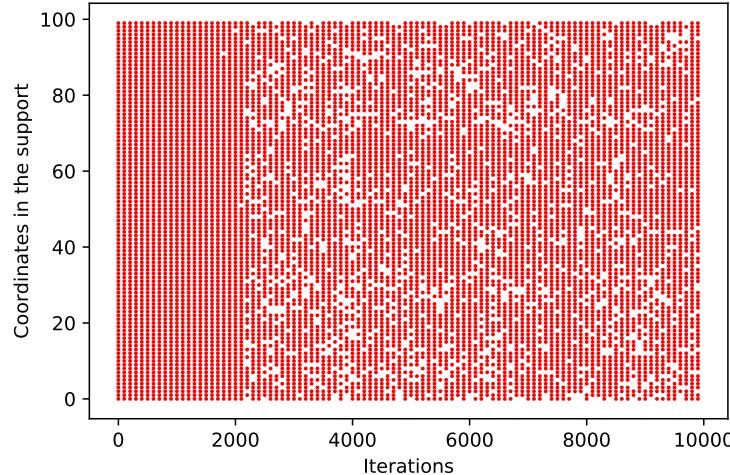
$$\min_{x \in \mathbb{R}^n} \frac{1}{m} \sum_{j=1}^m \log(1 + \exp(-y_j z_j^\top x)) + \lambda_1 \|x\|_1 + \frac{\lambda_2}{2} \|x\|_2^2$$

Why not SGD



Prox GD

Synthetic LASSO problem $\min \frac{1}{2} \|Ax - b\|_2^2 + \lambda_1 \|x\|_1$ for random generated matrix $A \in \mathbb{R}^{100 \times 100}$ and vector $b \in \mathbb{R}^{100}$ and hyperparameter λ_1 chosen to reach 15% of density (amount of non-zero coordinates) of the final solution.



Prox SGD (minibatch of size 10)



Non-degeneracy

Another way to define the non-degeneracy for the problem

$$\min_{x \in \mathbb{R}^n} f(x) + r(x)$$

is the following:

$$\nabla f(x^\star) \in \text{ri } \partial r(x^\star).$$

In case of ℓ_1 regularizer $r(x) = \lambda_1 \|x\|_1$ this can be written explicitly as

$$|\nabla f(x^\star)_{[j]}| < \lambda_1 \quad \text{for all } j \in \text{supp}(x^\star).$$